Sketchbook Upper Bounds

ABSTRACT

CCS CONCEPTS

• **Theory of computation → Bloom filters and hashing**; • **Information systems → Unidimensional range search**.

KEYWORDS

Range Filter, Dynamic Data Structure, Data Growth, Scalability.

ACM Reference Format:

. 2025. Sketchbook Upper Bounds. In *Proceedings of VLDB International Conference on Very Large Databases (2025) (VLDB '25).* ACM, New York, NY, USA, 3 pages. https://doi.org/XXXXXXX.XXXXXXX

1 THEORETICAL ANALYSIS AND RESULTS

We define $\text{PoE}_{\text{CMS}}(d, w, N)$ to denote the extra error we incur in CMS by not knowing the ℓ_1 norm of the stream in advance. Namely, we compare the error of Sketchbook with that of a CMS with d arrays of length w after it has processed a stream of ℓ_1 norm N.

this needs more work, because it does involve arguing about the probability of deviating from the expected value..

2 PROOF OF THEOREM FOR INSERTION-ONLY STREAMS

We will do the analysis assuming a general size function w^* and then discuss how the analysis proceeds for specific values of w^*. Note that both CS and CMS share the same general structure: they contain d counter arrays of size w.

Fix some size function w^*. Note that we perform doubling each time the function w^* doubles. In the CMS case, w^* is a function of N, the ℓ_1 norm of the stream. In the CS case, w^* is a function of the ℓ_2 norm of the stream. In either case, we can track the moments *in time* in which the value of the function doubles. Formally, instead of thinking of w^* as a function of the the ℓ_1 or ℓ_2 norms of the stream, we think of it as a function of time such that, at any point in time t, $w^*(t)$ is the number of counters the sketch uses at time t.

Now let T_0 be the first time the sketch is instantiated (i.e., this sketch has d counters and $w^*(T_0)$). We then define T_i for $i \geq 1$ to be the first moment in time when the length of each array doubles, i.e., $w^*(T_i) = 2 \cdot w^*(T_{i-1})$. Note that, in between times T_i and T_{i+1}, the length of each array stays the same (i.e., it is equal to $w^*(T_i)$). We call this period of time epoch i and denote by $\mathcal{S}_i$ the *active* sketch during this epoch (i.e., the one that has arrays of length $w^*(T_i)$). According to our design, updates to the stream during Epoch i are executed directly on $\mathcal{S}_i$. For the purposes of executing queries fast, we also aggregate all previous sketches into the active sketch.

Formally, let (x, Δ) be an update that occured in some epoch $0 \leq j < i$ before the current epoch. In the aggregation step, this update operation is replicated on all the counters in $\mathcal{S}_i$ that x could have hashed into. Namely, if x hashed to an index α in epoch j for one of the arrays, then, during epoch i, all of the 2^{j-1} counters $\alpha \circ \{0,1\}^{j-i}$ ingest the update (x, Δ), where $\circ$ denotes concatenation and $\{0,1\}^{j-i}$ denotes all binary strings of length $j-1$. From the perspective of fixed counter, what this means is that the counter must also ingest previous updates from elements that don't technically hash into it.

Formally, let α be a counter in epoch i, and assume for simplicity that we only have one array in the sketch (i.e., $d = 1$ and $h_1(x) = h(x)$). This counter ingests the following operations:

- updates (x, Δ) from epoch i and for which $h(x) = \alpha$. This event occurs with probability $1/w^*(T_i)$ since there are $w^*(T_i)$ counters in the array. For ease of analysis, we define $\mathcal{J}_{i,\alpha}(x)$ to be the indicator random variable that is 1 if this event occurs, and 0 otherwise. Then $\Pr\left[\mathcal{J}_{i,\alpha}(x) = 1\right] = 1/w^*(T_i)$.
- updates (x, Δ) from epoch $i-1$ that, during epoch $i-1$, hashed into a counter that matches the leftmost $\log w^*(T_i) - 1$ bits of α in the binary representation (meaning, their current $h(x)$ value differs from α in at most the last 1 bit of their binary representation). This event occurs with probability at most $2/w^*(T_i)$. We emphasize that this is merely an upper bound: the additional factor of 2 accounts for the mistakes that the aggregation step makes, i.e., processing updates from previous epoch that would not have hashed into α in the current epoch.
- in general, updates (x, Δ) from epoch $j < i$ and for which $h(x)$ differs from α only in the last $i - j$ bits of its binary representation. This event occurs with probability at most $2^{i-j}/w^*(T_i)$. Similarly as before, we define $\mathcal{J}_{j,\alpha}(x)$ to be the indicator random variable that is 1 if this event occurs, and 0 otherwise. Then $\Pr\left[\mathcal{J}_{j,\alpha}(x) = 1\right] \leq 2^{i-j}/w^*(T_i)$. Similarly as before, note that this is merely an upper bound: for x such that $h(x) = \alpha$, the $\Pr\left[\mathcal{J}_{j,\alpha}(x) = 1 | h(x) = \alpha\right] = 1$ no matter when the update to x was processed in the stream.

While updates from previous epochs have a higher chance of hashing into α than updates in epoch i, they are nevertheless bounded in number, i.e., during epoch i, we have $T_{i+1} - T_i$ insertions. In the following, we bound the effect this has on each counter. We first cluster all the insertions that occur during a given epoch into one big insertion per element. Formally, let $f_i(x)$ denote the total count of updates of the form (x, Δ) that occur in the stream during epoch i. Note that $\sum_x f_j(x) = T_{j+1} - T_j$ for all $j \geq 0$.

Bounds for CMS. In the case of CMS, each counter gathers the sum of frequencies of elements that hash into it. Namely, the random variable α can be expressed as:

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

VLDB '25, June 22–27, 2025, London, United Kingdom

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/18/06…$15.00
https://doi.org/XXXXXXX.XXXXXXX

$$\alpha = \sum_{j=0}^{i} \sum_{x} \mathcal{J}_{j,\alpha}(x) \cdot f_j(x) \ .$$

By linearity of expectation, we get that:

$$E[\alpha] = \sum_{j=0}^{i} \sum_{x} \frac{2^{i-j}}{w^*(T_i)} \cdot f_j(x) = \frac{1}{w^*(T_i)} \cdot \sum_{j=0}^{i} \sum_{x} 2^{i-j} \cdot f_j(x)$$
$$= \frac{1}{w^*(T_i)} \cdot \sum_{j=0}^{i} 2^{i-j} \sum_{x} f_j(x) = \frac{1}{w^*(T_i)} \cdot \sum_{j=0}^{i} 2^{i-j} \cdot (T_{j+1} - T_j) \ .$$

In other words, the increase in error we get because of aggregating the sketches is dominated by the term $\sum_{j=0}^{i} 2^{i-j} \cdot (T_{j+1} - T_j)$. This is where the specific size function w^* we employ has an effect, since it translates into bounds on T_i (and specifically their relative sizes). We now consider various special cases and discuss their respective behavior:

Setting $w^*(N) = N/\epsilon$ in CMS. In this case, we start with the original sketch $\mathcal{S}_0$ with a capacity for T_0 elements. The size function doubles whenever the number of elements doubles. In this case $T_{j+1} = 2 \cdot T_j$ for all $j \geq 0$ and, in general, $T_j = 2^j \cdot T_0$ for all $j \geq 0$. The additional error term then becomes:

$$\sum_{j=0}^{i} 2^{i-j} \cdot (T_{j+1} - T_j) = \sum_{j=0}^{i} 2^{i-j} \cdot (2^{j+1} \cdot T_0 - 2^j \cdot T_0)$$
$$= \sum_{j=0}^{i} 2^{i-j} \cdot 2^j \cdot T_0 = i \cdot 2^i \cdot T_0 \ .$$

Since the ℓ_1 norm of the stream doubles in each epoch, we have that $2^i \cdot T_0 = T_i$ and so, in expectation,

$$E[\alpha] = \frac{1}{w^*(T_i)} \cdot i \cdot T_i = i/\epsilon \ .$$

This is also an upper bound on the expected error we get for an element x. blah blah

Setting $w^*(N) = N^{1-c}/\epsilon$ in CMS, for some $0 < c < 1$. In this case, the size function doubles more rarely, whenever $T_{i+1} = 2^{1/(1-c)} \cdot T_i$. Since $0 < c < 1$, note that this implies that T_{i+1} is bigger than $2T_i$. This difference is why we avoid the extra dependency on i in the error term. Formally, we have that $T_i = 2^{i/(1-c)} \cdot T_0$ and thus:

$$\sum_{j=0}^{i} 2^{i-j} \cdot (T_{j+1} - T_j) = \sum_{j=0}^{i} 2^{i-j} \cdot (2^{(j+1)/(1-c)} \cdot T_0 - 2^{j/(1-c)} \cdot T_0)$$
$$= \sum_{j=0}^{i} 2^{i-j} \cdot 2^{j/(1-c)} \cdot T_0 \cdot (2^{1/(1-c)} - 1)$$
$$= 2^i \cdot T_0 \cdot (2^{1/(1-c)} - 1) \cdot \sum_{j=0}^{i} 2^{j/(1-c)-j} \ .$$

For simplicity, we define $2^{1/(1-c)-1} = a$ and observe that:

$$\sum_{j=0}^{i} 2^{j/(1-c)-j} = \sum_{j=0}^{i} a^j = (a^{i+1} - 1)/(a-1) \ .$$

We then get that:

$$\sum_{j=0}^{i} 2^{i-j} \cdot (T_{j+1} - T_j) = 2^i \cdot T_0 \cdot (2^{1/(1-c)} - 1) \cdot \sum_{j=0}^{i} 2^{j/(1-c)-j}$$
$$= 2^i \cdot T_0 \cdot (2a-1) \cdot (a^{i+1} - 1)/(a-1)$$
$$\leq 2^{i+1} \cdot T_0 \cdot (a^{i+1} - 1) \leq 2^{i+1} \cdot T_0 \cdot a^{i+1} \ .$$

At this point, it is worth mentioning that the expected error we are comparing against is $T_i/w^*(T_i)$, i.e., the error that we get by aggregating is close to the error we would get if we didn't expand and just allocated S_i as a sketch from the beginning. To this end, we first notice that $2a = 2^{1/(1-c)}$ and so

$$2^{i+1} \cdot T_0 \cdot a^{i+1} = 2^{(i+1)/(1-c)} \cdot T_0 = 2^{1/(1-c)} \cdot T_i \ .$$

Therefore, the expected error becomes:

$$E[\alpha] \leq \frac{1}{w^*(T_i)} \cdot \sum_{j=0}^{i} 2^{i-j} \cdot (T_{j+1} - T_j) \leq 2^{1/(1-c)} \cdot \frac{T_i}{w^*(T_i)} \ ,$$

as desired. We note that the above holds whenever $c < 1$. In particular, c does not need to be a constant.

Streams with deletions. to be filled

Bounds for CS. In the case of CS, we are interested in describing the behavior of the random variable $\hat{f}(z) = \alpha \cdot \sigma(z)$ for any query z, where $\sigma(z)$ is the associated sign function (for simplicity, we again assume that $d = 1$). Using the notation we introduced before, the random variable α can be expressed as:

$$\alpha = \sum_{j=0}^{i} \sum_{x} \mathcal{J}_{j,\alpha}(x) \cdot f_j(x) \cdot \sigma(x) = \sum_{x} C(x) \ ,$$

where $C(x) = \sum_{j=0}^{i} \mathcal{J}_{j,\alpha}(x) \cdot f_j(x) \cdot \sigma(x)$ denotes the contribution throughout the epoch of updates to x.

We begin by computing the expected value of the estimator $\hat{f}(z)$ as in the classical CS analysis that shows that in expectation, $\alpha \cdot \sigma(z)$ is an unbiased estimator, i.e. $E[\alpha \cdot \sigma(z)] = f(z)$. By linearity of expectation

$$\mathsf{E}[\hat{f}(z)] = \sum_{x \neq z} \mathsf{E}\left[C(x) \cdot \sigma(z)\right] + \mathsf{E}\left[C(z)\sigma(z)\right] = \sum_{j=0}^{i} f_j(z) = f(z) \ ,$$

where the last step uses the fact that $\mathsf{E}\left[C(x) \cdot \sigma(z)\right] = 0$ by construction and that $\Pr\left[\mathcal{J}_{j,\alpha}(z)\right] = 1$ since $h(z) = \alpha$. We will now bound the variance of this estimator, and in particular, we will be interested in how much bigger it becomes when compared to $\sqrt{F_i/w^*(T_i)}$, where F_i is the squared ℓ_2 norm of the stream at time T_i. Using a similar argument when expanding the variance of the estimator, we get that:

$$\mathsf{Var}\left[\hat{f}(z)\right] = \mathsf{E}\left[\left(\sum_{x} C(x)\sigma(z)\right)^2\right] - (f(z))^2$$
$$= \sum_{x} \mathsf{E}\left[(C(x))^2\right] + \sum_{x} \sum_{y \neq x} \mathsf{E}\left[C(x)C(y)\right] - (f(z))^2$$
$$= \sum_{x} \mathsf{E}\left[(C(x))^2\right] - (f(z))^2 \ ,$$

where we use the fact that $\mathsf{E}\left[\sigma(x)\cdot\sigma(y)\right] = 0$ for any $x \neq y$ to establish that $\mathsf{E}\left[C(x)C(y)\right] = 0$. We now bound $\mathsf{E}\left[(C(x))^2\right]$. Recall that we defined $C(x) = \sum_{j=0}^{i} \mathcal{J}_{j,\alpha}(x)\cdot f_j(x)\cdot\sigma(x)$. Then:

$$\begin{aligned}\mathsf{E}\left[(C(x))^2\right] &= \mathsf{E}\left[\left(\sum_{j=0}^{i}\mathcal{J}_{j,\alpha}(x)\cdot f_j(x)\right)^2\right]\\ &= \sum_{j=0}^{i}\mathsf{E}\left[\left(\mathcal{J}_{j,\alpha}(x)\cdot f_j(x)\right)^2\right] +\\ &+ 2\cdot\sum_{j=0}^{i-1}\sum_{k=j+1}^{i}\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\cdot f_j(x)\cdot\mathcal{J}_{k,\alpha}(x)\cdot f_k(x)\right]\\ &= \sum_{j=0}^{i} f_j(x)^2\cdot\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\right] +\\ &+ 2\cdot\sum_{j=0}^{i-1}\sum_{k=j+1}^{i} f_j(x)\cdot f_k(x)\cdot\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\cdot\mathcal{J}_{k,\alpha}(x)\right]\ .\end{aligned}$$

For bounding the $\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\right]$, we can use the fact that $\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\right] \leq 2^{i-j}/w^*(T_i)$. For $\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\cdot\mathcal{J}_{k,\alpha}(x)\right]$, we note that the product of the two indicator random variables is 1 only when both of them are 1. Having said that, the two random variables are not independent: if $\mathcal{J}_{k,\alpha}(x) = 1$, then $\mathcal{J}_{j,\alpha}(x) = 1$ for any $k \geq j$. Namely, if during epoch k, x hashed into a counter that was a prefix of α, then the same would have happened during epoch $j \leq k$. We get that:

$$\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\cdot\mathcal{J}_{k,\alpha}(x)\right] = \mathsf{E}\left[\mathcal{J}_{k,\alpha}(x)\right]\ .$$

The whole sums becomes:

$$\begin{aligned}\sum_{j=0}^{i-1}\sum_{k=j+1}^{i} f_j(x)\cdot f_k(x)\cdot\mathsf{E}\left[\mathcal{J}_{k,\alpha}(x)\right] &= \sum_{k=1}^{i}\sum_{j=0}^{k-1} f_j(x)\cdot f_k(x)\cdot\mathsf{E}\left[\mathcal{J}_{k,\alpha}(x)\right]\\ &= \sum_{k=1}^{i} f_k(x)\cdot\mathsf{E}\left[\mathcal{J}_{k,\alpha}(x)\right]\cdot\sum_{j=0}^{k-1} f_j(x)\\ &= \sum_{k=1}^{i} f_k(x)\cdot\mathsf{E}\left[\mathcal{J}_{k,\alpha}(x)\right]\cdot f_{\leq k-1}(x)\end{aligned}$$

where, for simplicity, we denoted $f_{\leq i}(x) = \sum_{j=0}^{i} f_i(x)$ the sum of frequencies accumulated by x through epochs 0 to i and set $f_{\leq -1}(x) = 0$. Then:

$$\begin{aligned}\mathsf{E}\left[(C(x))^2\right] &= \sum_{j=0}^{i}\left(f_j(x)^2 + 2f_j(x)\cdot f_{\leq j-1}(x)\right)\cdot\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\right]\\ &= \sum_{j=0}^{i}\left(f_{\leq j}(x)^2 - f_{\leq j-1}(x)^2\right)\cdot\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\right]\ .\end{aligned}$$

While the term $f_{\leq j}(x)^2 - f_{\leq j-1}(x)^2$ might seem contrited, we note that $\sum_x f_{\leq j}(x)^2 - f_{\leq j-1}(x)^2 = F_j - F_{j-1}$ denotes the increase in the squared ℓ_2 norm of the stream vector as we move from epoch $j-1$ to epoch j. We thus get that:

$$\begin{aligned}\mathsf{Var}\left[\hat{f}(z)\right] &= \sum_x \mathsf{E}\left[(C(x))^2\right] - (f(z))^2\\ &= \sum_x\sum_{j=0}^{i}\left(f_{\leq j}(x)^2 - f_{\leq j-1}(x)^2\right)\cdot\mathsf{E}\left[\mathcal{J}_{j,\alpha}(x)\right] - (f(z))^2\\ &\leq \sum_{j=0}^{i}\left(F_j - F_{j-1}\right)\cdot\frac{2^{i-j}}{w^*(T_i)}\ .\end{aligned}$$

At this point, it becomes perhaps more natural to understand why we defined the size function to depend on the squared ℓ_2 norm of the stream vector. We now discuss different bounds we get on the variance depending on the size function.

Setting $w^*(F) = F/\epsilon$ in CS. In this case, we start with the original sketch $\mathcal{S}_0$ that allows streams for which their squared ℓ_2 norm is at most some F_0. The size function doubles whenever the squared ℓ_2 norm doubles. In this case $F_{j+1} = 2\cdot F_j$ for all $j \geq 0$ and, in general, $F_j = 2^j\cdot F_0$ for all $j \geq 0$. In this case, we get that:

$$\begin{aligned}\mathsf{Var}\left[\hat{f}(z)\right] &\leq \sum_{j=0}^{i}\left(F_j - F_{j-1}\right)\cdot\frac{2^{i-j}}{w^*(T_i)}\\ &\leq \frac{1}{w^*(T_i)}\cdot\left(F_0\cdot 2^i + \sum_{j=1}^{i}\left(2^j - 2^{j-1}\right)\cdot F_0\cdot 2^{i-j}\right)\\ &\leq \frac{1}{w^*(T_i)}\cdot\left(F_0\cdot 2^i + \sum_{j=1}^{i} 2^{j-1}\cdot F_0\cdot 2^{i-j}\right)\\ &\leq \frac{1}{w^*(T_i)}\cdot\left(F_0\cdot 2^i + \sum_{j=1}^{i} F_0\cdot 2^{i-1}\right)\\ &\leq \frac{1}{w^*(T_i)}\cdot\left(F_i + (i-1)\cdot F_i/2\right) = \frac{F_i}{w^*(T_i)}\cdot\frac{i+1}{2}\ .\end{aligned}$$

Received 16 April 2024; revised 4 June 2024; accepted 8 August 2024